


Propositional Logic and AI



The Wumpus World

- Performance measure
 - gold +1000, death -1000
 - -1 per step, -10 for using the arrow
- Environment
 - Squares adjacent to wumpus are smelly
 - Squares adjacent to pit are breezy
 - Glitter iff gold is in the same square
 - Shooting kills wumpus if you are facing it
 - Shooting uses up the only arrow
 - Grabbing picks up gold if in same square
 - Releasing drops the gold in same square
 - Walking into a wall makes the agent perceive a bump
 - When the wumpus is killed, it emits a scream that is heard throughout the cave
- Sensors: Stench, Breeze, Glitter, Bump, Scream
- Actuators: turn left, turn right, move forward, Grab, Release, Shoot



Exploring a wumpus world

OK			
OK <div>A</div>	OK		

Logic for reasoning


- Goal: Deduce new facts (α) using:
 - The rules about the world
 - Information we gather through perception
- } KB
- **Entailment** means that one thing **follows from** another:

$$KB \models \alpha$$

Last night I ate either spaghetti or pizza
I did not eat pizza last night \models Last night I ate spaghetti

Model

- Assignment of a truth value – true or false – to every atomic sentence



M_1	$S = \text{False}$ $P = \text{False}$
M_2	$S = \text{False}$ $P = \text{True}$
M_3	$S = \text{True}$ $P = \text{False}$
M_4	$S = \text{True}$ $P = \text{True}$

A model m is a model of KB iff it is a model of all sentences in KB, that is, all sentences in KB are true in m

Satisfiability of a KB

A KB is **satisfiable** iff it admits at least one model; otherwise it is **unsatisfiable**


KB1 = $\{P, \neg Q \wedge R\}$ is _____

KB2 = $\{\neg P \vee P\}$ is _____

KB3 = $\{P, \neg P\}$ is _____

Logical Entailment

- KB : set of sentences
- α : arbitrary sentence
- KB **entails** α – written $\text{KB} \models \alpha$ – iff every model of KB is also a model of α
- Alternatively, $\text{KB} \models \alpha$ iff
 - $\{\text{KB}, \neg\alpha\}$ is unsatisfiable
 - $\text{KB} \Rightarrow \alpha$ is valid



KB

S

α

M_1

$S = \text{False}$
 $P = \text{False}$

M_2

$S = \text{False}$
 $P = \text{True}$

M_3

$S = \text{True}$
 $P = \text{False}$

M_4


$S = \text{True}$
 $P = \text{True}$

Entailment in the wumpus world


Situation after detecting
nothing in [1,1], moving
right, breeze in [2,1]

Consider possible models for
KB assuming only pits

3 Boolean choices \Rightarrow 8
possible models




Wumpus models




- KB = wumpus-world rules + 2 observations
- α_1 = "[1,2] is safe"

Wumpus models




- KB = wumpus-world rules + 2 observations
- α_1 = "[1,2] is safe"

Wumpus models




- KB = wumpus-world rules + observations
- α_1 = "[1,2] is safe"

Wumpus models




- KB = wumpus-world rules + 2 observations
- α_2 = "[2,2] is safe"

Wumpus models



- KB = wumpus-world rules + 2 observations
- α_2 = "[2,2] is safe"

Wumpus models



- KB = wumpus-world rules + observations
- α_2 = "[2,2] is safe"

Propositional logic: Syntax

- Propositional logic is the simplest logic – illustrates basic ideas
- The proposition symbols P_1 , P_2 etc are sentences
 - If S is a sentence, $\neg S$ is a sentence (**negation**)
 - If S_1 and S_2 are sentences, $S_1 \wedge S_2$ is a sentence (**conjunction**)
 - If S_1 and S_2 are sentences, $S_1 \vee S_2$ is a sentence (**disjunction**)
 - If S_1 and S_2 are sentences, $S_1 \Rightarrow S_2$ is a sentence (**implication**)
 - If S_1 and S_2 are sentences, $S_1 \Leftrightarrow S_2$ is a sentence (**biconditional**)

Propositional logic: Semantics

Each model specifies true/false for each proposition symbol

E.g. $P_{1,2}$ $P_{2,2}$ $P_{3,1}$
false true false

With these symbols, 8 possible models, can be enumerated automatically.

Rules for evaluating truth with respect to a model m :

$\neg S$	is true iff	S is false
$S_1 \wedge S_2$	is true iff	S_1 is true and S_2 is true
$S_1 \vee S_2$	is true iff	S_1 is true or S_2 is true
$S_1 \Rightarrow S_2$	is true iff	S_1 is false or S_2 is true
$S_1 \Leftrightarrow S_2$	is true iff	$S_1 \Rightarrow S_2$ is true and $S_2 \Rightarrow S_1$ is true

Simple recursive process evaluates an arbitrary sentence, e.g.,

$$\neg P_{1,2} \wedge (P_{2,2} \vee P_{3,1}) = \text{true} \wedge (\text{true} \vee \text{false}) = \text{true} \wedge \text{true} = \text{true}$$

Inference


- Just because a KB entails a sentence doesn't mean we can find (infer) it
- Inference is the process of generating sentences entailed by the KB



Inference Rule

- An **inference rule** $\{\xi, \psi\} \vdash \varphi$ consists of 2 sentence patterns ξ and ψ called the **conditions** and one sentence pattern φ called the **conclusion**
- If ξ and ψ match two sentences of KB then the corresponding φ can be inferred according to the rule

Example: Modus Ponens



Battery-OK \wedge Bulbs-OK \Rightarrow Headlights-Work ($\alpha \Rightarrow \beta$)

Battery-OK \wedge Starter-OK \wedge \neg Empty-Gas-Tank \Rightarrow Engine-Starts

Engine-Starts \wedge \neg Flat-Tire \Rightarrow Car-OK

Battery-OK \wedge Bulbs-OK (α)

Can be used as inference rules...

$(\alpha \wedge \beta) \equiv (\beta \wedge \alpha)$	commutativity of \wedge
$(\alpha \vee \beta) \equiv (\beta \vee \alpha)$	commutativity of \vee
$((\alpha \wedge \beta) \wedge \gamma) \equiv (\alpha \wedge (\beta \wedge \gamma))$	associativity of \wedge
$((\alpha \vee \beta) \vee \gamma) \equiv (\alpha \vee (\beta \vee \gamma))$	associativity of \vee
$\neg(\neg\alpha) \equiv \alpha$	double-negation elimination
$(\alpha \Rightarrow \beta) \equiv (\neg\beta \Rightarrow \neg\alpha)$	contraposition
$(\alpha \Rightarrow \beta) \equiv (\neg\alpha \vee \beta)$	implication elimination
$(\alpha \Leftrightarrow \beta) \equiv ((\alpha \Rightarrow \beta) \wedge (\beta \Rightarrow \alpha))$	biconditional elimination
$\neg(\alpha \wedge \beta) \equiv (\neg\alpha \vee \neg\beta)$	de Morgan
$\neg(\alpha \vee \beta) \equiv (\neg\alpha \wedge \neg\beta)$	de Morgan
$(\alpha \wedge (\beta \vee \gamma)) \equiv ((\alpha \wedge \beta) \vee (\alpha \wedge \gamma))$	distributivity of \wedge over \vee
$(\alpha \vee (\beta \wedge \gamma)) \equiv ((\alpha \vee \beta) \wedge (\alpha \vee \gamma))$	distributivity of \vee over \wedge

Figure 7.11 Standard logical equivalences. The symbols α , β , and γ stand for arbitrary sentences of propositional logic.

\Rightarrow Connective symbol (implication)

\models Logical entailment

$KB \models \alpha$ iff $KB \Rightarrow \alpha$ is valid

\vdash Inference

Soundness

- An inference rule is **sound** if it generates only entailed sentences
- All inference rules previously given are sound, e.g.:
modus ponens: $\{\alpha \Rightarrow \beta, \alpha\} \vdash \beta$
- Is the following rule sound?
 $\{\alpha \Rightarrow \beta, \beta\} \vdash \alpha$

Completeness

- A set of inference rules is **complete** if every entailed sentences can be obtained by applying some finite succession of these rules
- Modus ponens *alone* is not complete, e.g.:
from $A \Rightarrow B$ and $\neg B$, we cannot get $\neg A$

Proof

The **proof** of a sentence α from a set of sentences KB is the derivation of α by applying a series of sound (legal) inference rules

Proof

The **proof** of a sentence α from a set of sentences KB is the derivation of α by applying a series of sound inference rules

- | | | |
|----|----|---|
| 1 | 1. | Battery-OK \wedge Bulbs-OK \Rightarrow Headlights-Work |
| 2 | 2. | Battery-OK \wedge Starter-OK \wedge \neg Empty-Gas-Tank \Rightarrow Engine-Starts |
| 3 | 3. | Engine-Starts \wedge \neg Flat-Tire \Rightarrow Car-OK |
| 4 | 4. | Headlights-Work |
| 5 | 5. | Battery-OK |
| 6 | 6. | Starter-OK |
| 7 | 7. | \neg Empty-Gas-Tank |
| 8 | 8. | \neg Car-OK |
| 9 | | |
| 10 | | |
| 11 | | |
| 12 | | |
| 13 | | |

Proof

The **proof** of a sentence α from a set of sentences KB is the derivation of α by applying a series of sound inference rules

- | | | | |
|----|-----|---|----------------------|
| 1 | 1. | Battery-OK \wedge Bulbs-OK \Rightarrow Headlights-Work | |
| 2 | 2. | Battery-OK \wedge Starter-OK \wedge \neg Empty-Gas-Tank \Rightarrow Engine-Starts | |
| 3 | 3. | Engine-Starts \wedge \neg Flat-Tire \Rightarrow Car-OK | |
| 4 | 4. | Headlights-Work | |
| 5 | 5. | Battery-OK | |
| 6 | 6. | Starter-OK | |
| 7 | 7. | \neg Empty-Gas-Tank | |
| 8 | 8. | \neg Car-OK | |
| 9 | 9. | Battery-OK \wedge Starter-OK | $\leftarrow (5+6)$ |
| 10 | 10. | Battery-OK \wedge Starter-OK \wedge \neg Empty-Gas-Tank | $\leftarrow (9+7)$ |
| 11 | 11. | Engine-Starts | $\leftarrow (2+10)$ |
| 12 | 12. | Engine-Starts \Rightarrow Flat-Tire | $\leftarrow (3+8)$ |
| 13 | 13. | Flat-Tire | $\leftarrow (11+12)$ |

Inference Problem

- **Given:**
 - KB: a set of sentence
 - α : a sentence
- **Answer:**
 - $\text{KB} \models \alpha$?

We require an automatic, sound and complete inference method

Inference by enumeration

- We can enumerate all possible models and test whether every model of KB is also a model of α

Wumpus world sentences

Let $P_{i,j}$ be true if there is a pit in $[i, j]$

Let $B_{i,j}$ be true if there is a breeze in $[i, j]$

KB:

$\neg P_{1,1}$

$\neg B_{1,1}$

$B_{2,1}$

Rules of the environment:

"Pits cause breezes in adjacent squares"

Truth tables for inference

$B_{1,1}$	$B_{2,1}$	$P_{1,1}$	$P_{1,2}$	$P_{2,1}$	$P_{2,2}$	$P_{3,1}$	KB	α_1
false	false	false	false	false	false	false	false	true
false	false	false	false	false	false	true	false	true
\vdots	\vdots	\vdots	\vdots	\vdots	\vdots	\vdots	\vdots	\vdots
false	true	false	false	false	false	false	false	true
false	true	false	false	false	false	true	<u>true</u>	<u>true</u>
false	true	false	false	false	true	false	<u>true</u>	<u>true</u>
false	true	false	false	false	true	true	<u>true</u>	<u>true</u>
false	true	false	false	true	false	false	false	true
\vdots	\vdots	\vdots	\vdots	\vdots	\vdots	\vdots	\vdots	\vdots
true	true	true	true	true	true	true	false	false

$\neg P_{1,1}$

$\neg B_{1,1}$

$B_{2,1}$

$B_{1,1} \Leftrightarrow (P_{1,2} \vee P_{2,1})$

$B_{2,1} \Leftrightarrow (P_{1,1} \vee P_{2,2} \vee P_{3,1})$

$\alpha_1 \neg P_{1,2}$

Inference by enumeration

- How efficient is this? How much time/space does it take?

Proof methods

- Proof methods divide into (roughly) two kinds:
 - Application of inference rules
 - Legitimate (sound) generation of new sentences from old
 - **Proof** = a sequence of inference rule applications
Can use inference rules as operators in a standard search algorithm
 - Typically require transformation of sentences into a **normal form**
 - Model checking
 - truth table enumeration (always exponential in n)
 - improved backtracking, e.g., Davis--Putnam-Logemann-Loveland (DPLL)
 - heuristic search in model space (sound but incomplete)
e.g., min-conflicts-like hill-climbing algorithms

Resolution Inference Rule

Conjunctive Normal Form (CNF)

conjunction of disjunctions of literals clauses

E.g., $(A \vee \neg B) \wedge (B \vee \neg C \vee \neg D)$


- Resolution inference rule (for CNF):

$$\frac{\begin{array}{c} \ell_i \vee \dots \vee \ell_k, \\ m_1 \vee \dots \vee m_n \end{array}}{\ell_i \vee \dots \vee \ell_{i-1} \vee \ell_{i+1} \vee \dots \vee \ell_k \vee m_1 \vee \dots \vee m_{j-1} \vee m_{j+1} \vee \dots \vee m_n}$$

where ℓ_i and m_j are complementary literals.

$$\text{E.g., } \frac{P_{1,3} \vee P_{2,2}, \quad \neg P_{2,2}}{P_{1,3}}$$

- Resolution is sound and complete for propositional logic



Conjunctive Normal Form

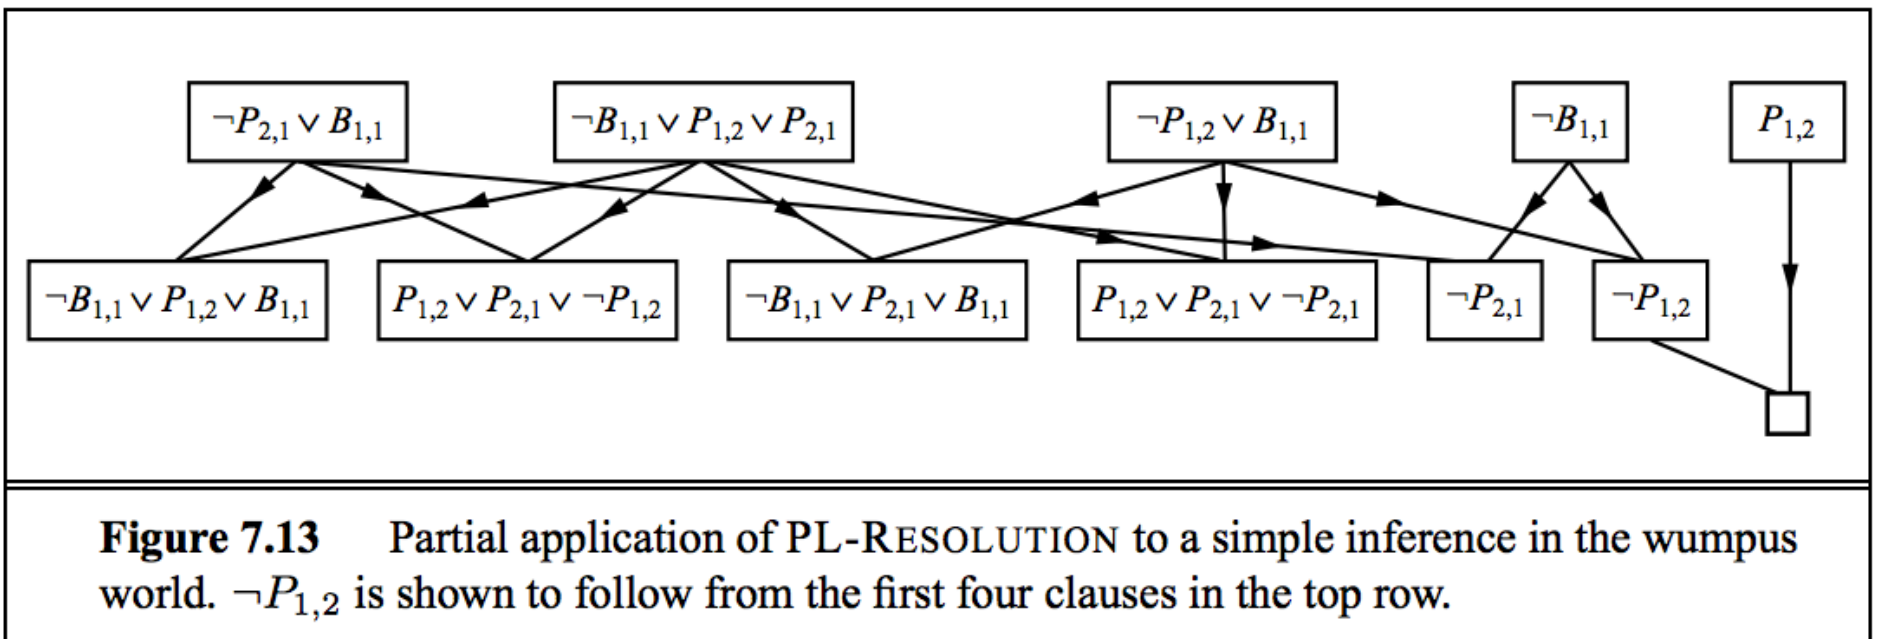
- Resolution rule only applies to disjunctions of literals
- How could it possibly be complete?
- EVERY sentence in prop logic is equivalent to a conjunction of clauses (disjunction of literals)
- This equivalent form is called CNF
- We have a standard procedure for converting to CNF

Prove that $KB \models \alpha$ by proving that $KB \wedge \neg \alpha$ is unsatisfiable

```
function PL-RESOLUTION( $KB, \alpha$ ) returns true or false
  inputs:  $KB$ , the knowledge base, a sentence in propositional logic
            $\alpha$ , the query, a sentence in propositional logic

   $clauses \leftarrow$  the set of clauses in the CNF representation of  $KB \wedge \neg \alpha$ 
   $new \leftarrow \{ \}$ 
  loop do
    for each  $C_i, C_j$  in  $clauses$  do
       $resolvents \leftarrow$  PL-RESOLVE( $C_i, C_j$ )
      if  $resolvents$  contains the empty clause then return true
       $new \leftarrow new \cup resolvents$ 
  if  $new \subseteq clauses$  then return false
   $clauses \leftarrow clauses \cup new$ 
```

Figure 7.12 A simple resolution algorithm for propositional logic. The function PL-RESOLVE returns the set of all possible clauses obtained by resolving its two inputs.



Forward and backward chaining

- **Horn Form** (restricted)

KB = conjunction of Horn clauses

- Horn clause =
 - proposition symbol; or
 - (conjunction of symbols) \Rightarrow symbol
- E.g., $C \wedge (B \Rightarrow A) \wedge (C \wedge D \Rightarrow B)$

- **Modus Ponens** (for Horn Form): complete for Horn KBs


$$\begin{array}{ccc} \alpha_1, \dots, \alpha_n, & & \alpha_1 \wedge \dots \wedge \alpha_n \Rightarrow \beta \\ & \beta & \end{array}$$

- Can be used with **forward chaining** or **backward chaining**.
- These algorithms are very natural and run in **linear** time


Forward chaining

- Idea: fire any rule whose premises are satisfied in the *KB*,
 - add its conclusion to the *KB*, until query is found


Forward chaining example




Forward chaining example




Forward chaining example




Forward chaining example




Forward chaining example




Forward chaining example



Forward chaining example



Forward chaining example



Backward chaining

Idea: work backwards from the query q :

to prove q by BC,

check if q is known already, or


prove by BC all premises of some rule concluding q

Avoid loops: check if new subgoal is already on the goal stack


Avoid repeated work: check if new subgoal

1. has already been proved true, or
2. has already failed


Backward chaining example




Backward chaining example




Backward chaining example




Backward chaining example




Backward chaining example




Backward chaining example




Backward chaining example




Backward chaining example



Backward chaining example



Backward chaining example



Forward vs. backward chaining

- FC is **data-driven**, automatic, unconscious processing,
 - e.g., object recognition, routine decisions
- May do lots of work that is irrelevant to the goal
- BC is **goal-driven**, appropriate for problem-solving,
 - e.g., Where are my keys? How do I get into a PhD program?
- Complexity of BC can be **much less** than linear in size of KB